



# CLASS 8

## Final Project Topics

# Final Project Topics

- We will try to cover the “big” features first
- These are not intended to give you the full implementation details
  - Think “high-level overview”
  - You *will* have to do your own research
  - You’re focusing on these 1 or 2 things for multiple weeks, so make sure you’re interested in them
- Don’t wait for us to teach you – start your research now!
  - The staff has pretty good coverage, but even we don’t know some of this stuff
  - Teach us!



# CLASS 8

## Networking

Networking

# NETWORKING STRATEGIES

# The Illusion

- All players are playing in real-time on the same machine
- But of course this isn't possible
- We need to emulate this as much as possible



# The Illusion

- What the player should see:
  - Consistent game state
  - Responsive controls
  - Difficult to cheat
- Things working against us:
  - Game state > bandwidth
  - Variable or high latency
  - Antagonistic users



# Send the Entire World!

- Players take turns modifying the game world and pass it back and forth
- Works alright for turn-based games
- ...but usually it's bad
  - RTS: there are a million units
  - FPS: there are a million players
  - Fighter: timing is crucial



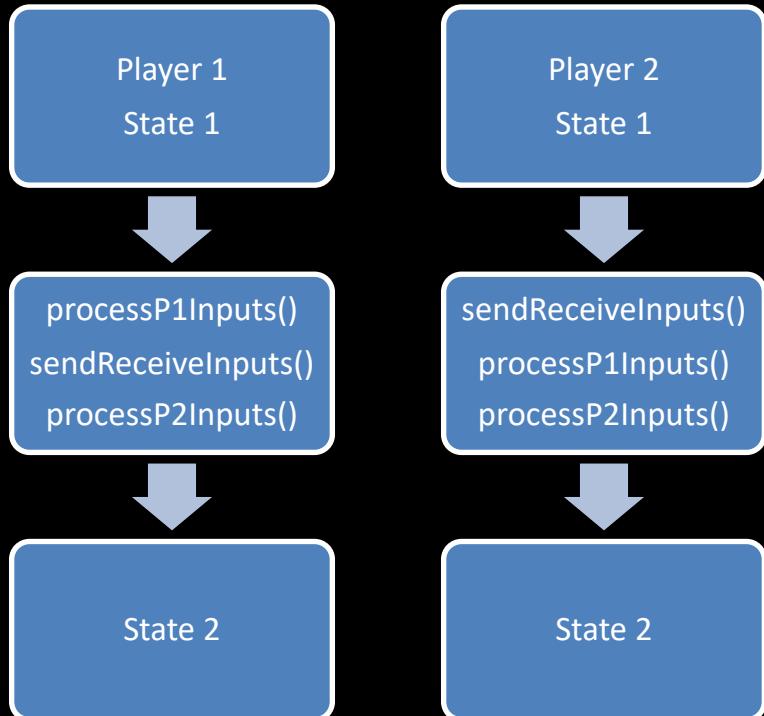
# Modeling the World

- If we're sending everything, we're modeling the world as a uniform chunk
  - But it really isn't!
  - Composed of entities, only some of which need input from a player
- We need a better model to solve these problems



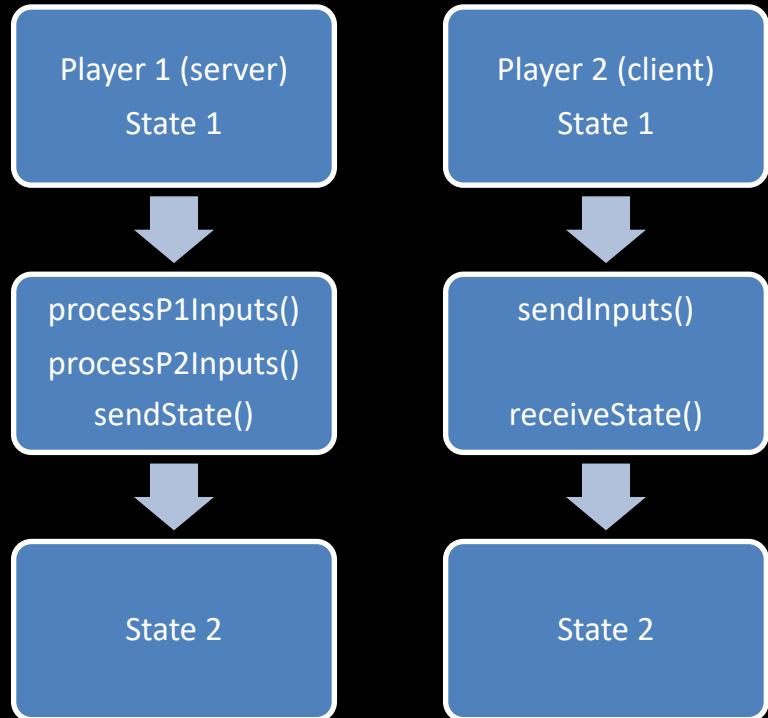
# Send Commands

- Model the world as local and shared data
  - Share player information, powerups, etc
  - Don't need to share static level data
- Each player sends the other all actions that alter shared game world
- “Deterministic P2P Lockstep”
- Problem: everything must evaluate the same
  - Or else there are desyncs
- Problem: have to wait for all the other players' commands
  - So everyone is limited by laggiest player



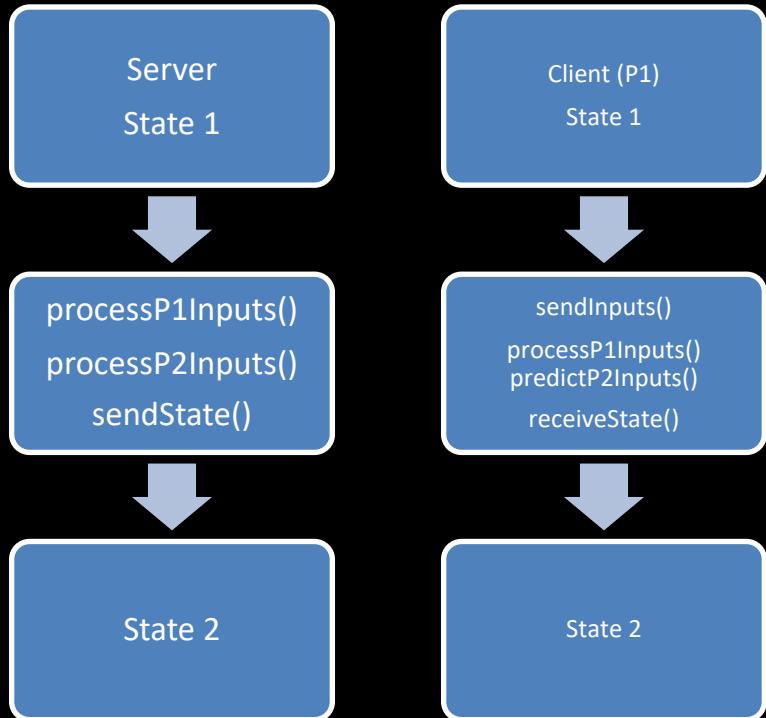
# Client-Server Model

- One player is the authoritative server
  - Now we don't have to wait for slow players, just the server
- Other player is a “dumb terminal”
  - Sends all input to server
  - Server updates the world and sends it back
- Problem: client has to wait for server to respond to perform even basic actions



# Client-side Prediction & Rollback

- Client responds to player input immediately and predicts new game state
  - This might include predicting other player's inputs
- When the server sends back the authoritative game state, incorrectly predicted client state is overwritten

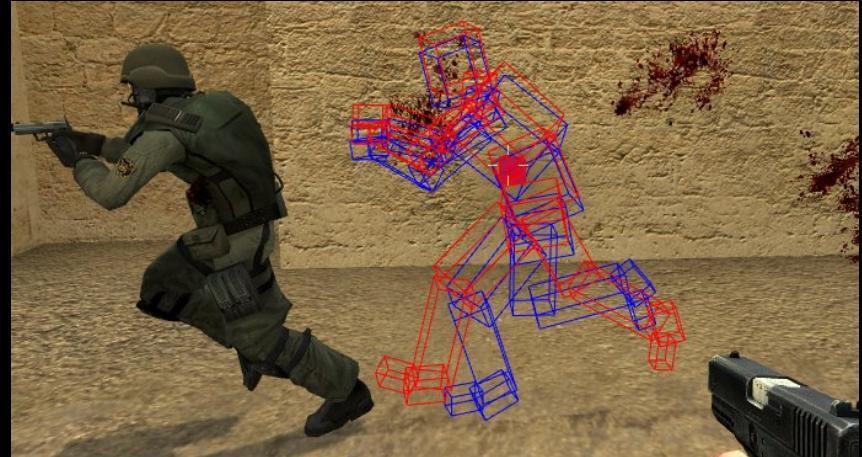


# Client-side Prediction & Rollback

- But we just received a state from the server that was 100ms in the past!
- We can't just replace our local state or else our local state would also be 100ms in the past
  - We lose player input that happened in those 100ms
- Client has to roll back the world and re-apply input that happened since the last known good state
  - This includes predicting new inputs for other players

# What about the server?

- Without rollback:
  - In an FPS, would need to lead shots because the server won't register shot until after delay
- With rollback:
  - You could be shot after you think you've taken cover



# Masking the Timewarp

- Problem: laggy players re-apply lots of inputs during rollback
- Solution: if the server usually sends states from 100ms ago, apply inputs 100ms late on the client
  - Still send inputs to the server immediately
- Turns a jumpy experience into a smooth, only slightly slow one
  - Very useful if relative timing of commands is important



# Edge Cases

- What if...
  - The client disconnects
  - The server dies
  - The client goes insane and sends gibberish
  - The client loses internet for 30 seconds
  - The client is malicious
  - The client changes IP address
- Handling errors well is vital to player experience

# Elegant Disconnects

- Handle and respond to IO exceptions
  - Don't just dump a stack trace
- Display informative status messages
- Send heartbeat packets every few seconds
  - Then respond if server/client hasn't received a heartbeat in a while
- Never let the game continue to run in an unrecoverable state!



# Resources

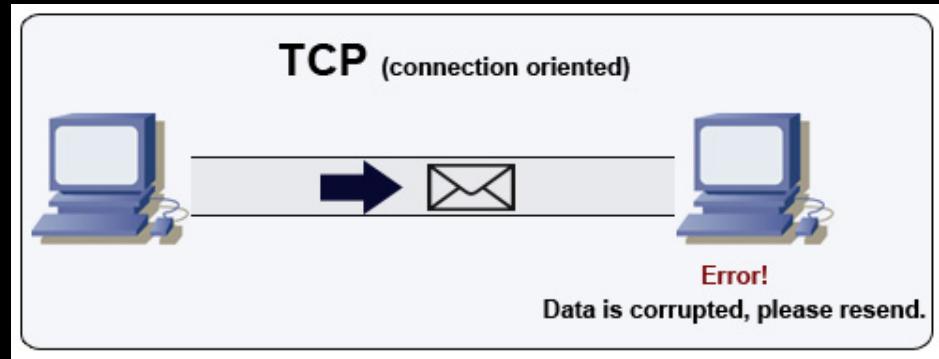
- For a more in-depth discussion about networking models and the concepts of client-side prediction and rollback
  - See [this link](#)
- Lots of other resources online

Networking

# IMPLEMENTATION

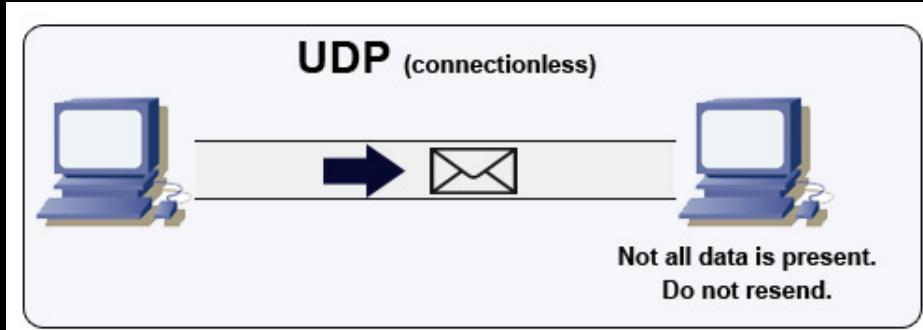
# TCP: Transmission Control Protocol

- Abstracts over IP
- All packets are guaranteed to be received and in the correct order
- Good for sending important, permanent data (websites, databases, etc)



# UDP: User Datagram Protocol

- A very thin shell around IP
- Much faster than TCP, but no guarantees about reception or order
- Good for information where only the most recent state matters (streaming, etc)



# TCP vs UDP

- (Very) generally: action games use UDP and turn-based games use TCP
  - World state updates can be lost without worry, commands not so much
- Can potentially combine them
  - TCP sends important data, UDP sends timely data
- Best choice varies by project
  - (for naïve version, TCP is fine)

# C sockets

- So much more than we want to cover in class
- Pros:
  - Full control over network throughput
  - Worth 5+ points
  - You will learn a lot
- Cons:
  - Oh so much more complicated
  - Will require multithreading, synchronization, and an incredibly well thought out design
- Start with CS033's snowcast project
  - <http://cs.brown.edu/courses/csci1972/handout/snowcast.pdf>

# C++ QSockets

- Qt has QTcpSocket and QUdpSocket!
- Pros:
  - Far easier to set up than standard sockets
  - Convenient blocking, non-blocking IO calls
- Cons:
  - Still sending/reading bytes
  - Still need multithreading, synchronization, and a good design
  - Ton of error checking required
- Better, but still not perfect. So...?

Networking

**RAKNET**

# The RakNet library

- Open-source games networking library
  - Recently bought by Oculus!
  - Plugin-style
- Used by some \*really\* legit engines
  - Unity, Havok, Minecraft
- Find it here: <http://www.jenkinssoftware.com/>

# The basics

- Basic client-server or P2P connections
  - Read and write threads made for you!
- BitStreams that can serialize:
  - Primitives (char, int, long, etc...)
  - !!! Structs !!!
- Basic packet objects with metadata

# Isn't that everything...?

- Now sending data is easy...great!
- Still have to...
  - Pick what to send
  - Pick when to send
  - Interpret what is sent
- What if I have 1000 entities in my world?
  - Entire world may be too much data...
  - Need some complex ID system
- Gee, it would be great if...

# RakNet does that too!

- Introducing ReplicaManager3!
  - Networked entities inherit from Replica3
  - RakNet gives you callbacks for all serialization events
    - onConstruct
    - onDestruct
    - onSerialize (each tick)
- \*A LOT\* of setup required, but works amazingly well
- Probably better to extend it a bit for simpler callbacks

# Using ReplicaManager3

- Have entities override some “NetworkedEntity” class that does most of the setup
  - Most of it is the same for every object
- Determine where entities are made and destroyed
  - Client or server side?
- Override serialization methods
  - Feed stuff in/out of a BitStream in the right order!
- Register created/destroyed entities with RakNet
- ???
- Profit!

# Sounds great!

- Since you decide what's serialized, you can avoid sending things other clients don't care about...
  - But what about things that don't change often?
- VariableDeltaSerializer!
  - Only sends data that hasn't changed since last tick!
  - More work now, less network throughput
- Space-time tradeoff probably worth it...

# I'm Sold!

- Lobby system
- “Fully connected mesh” host determination system
- Authentication protocols
- Team management
- !!! Voice chat !!!
- SQLite3 databases
- And so much more...

# In conclusion...

- RakNet is a beast
- Pros:
  - Handles the nitty-gritty threads/sockets details for you
  - RM3 really simplifies the design process
  - Lets you focus more on engine design
  - Fast.
- Cons:
  - A lot of setup required for RM3
  - You won't learn as much ☹

Networking - RakNet

**QUESTIONS?**



# CLASS 8

## Advanced Graphics

Advanced Graphics

# PARTICLES

# Particles

- What is a particle?
  - A particle is a tiny entity that is used in massive quantities to create a visually pleasing effect
  - Used to model effects like fire, liquids, hair, etc.
  - Conceptually old—first paper published in 1983

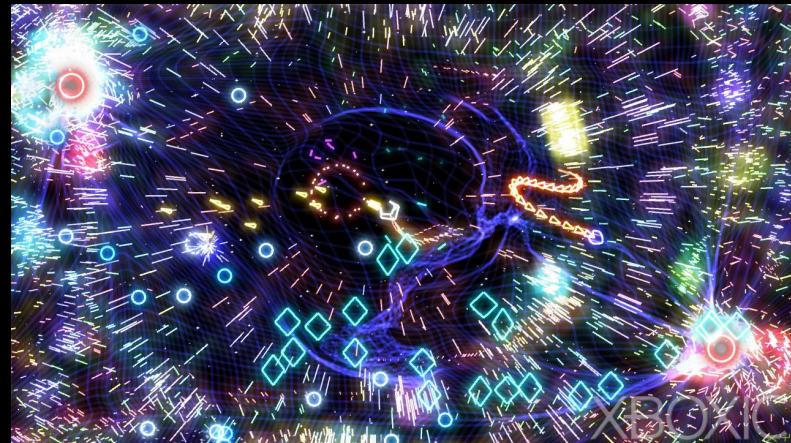
# Particles

- What makes particles look good?
- Fade out or get smaller linearly over their lifespan
  - Once the particle is completely gone or transparent, it can be removed from the world
- Adding some kind of randomness to how they move
  - Starting position, velocity, acceleration, color, size, shape



# Particles

- Particles are great
- But they are very slow if not done correctly
- Things that make them slow:
  - It's a lot of information to tick
  - It's a lot of information to draw
  - There are way too many of them to consider doing collision detection against each other



# Particles

- What shape should my particles be?
- Sphere?
  - Is limited if you want transparent particles
  - Too many vertices
- Quad!
  - Use a texture like this one
  - Rotate the quad to always face the camera
  - This texture has a black background and no alpha information
    - Use graphics->enableBlendTest(Graphics::BLEND\_FUNC::ADD)
    - This says take all of the background, and add the color of this particle on top of it
    - Particles that are denser will appear brighter



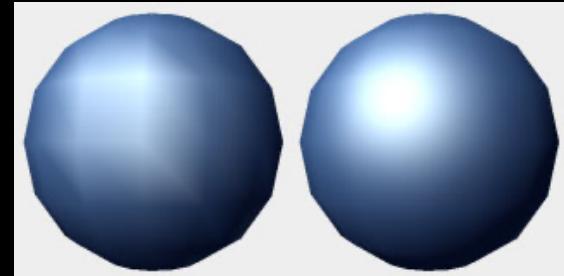
# Particle Optimizations

- Reduce the amount of information in your particles
  - Vector3 position, Vector3 velocity
  - maybe some noise values to make them scatter
  - Less information to tick
- Don't make your particles GameObjects
  - Keep them in a separate list so that you can tick and draw them all at once
  - Binding the particle texture once and then drawing all your particles without unbinding and rebinding the texture is a HUGE improvement
- Use GL\_TRIANGLE\_STRIP instead of GL\_TRIANGLES triangle strips to draw
  - Less data needs to be sent to the GPU with triangle strips
  - A particle is just a quad, so you only have 4 vertices vs 6 vertices
- Don't collide them with each other
  - If you must have your particles collide, have them collide only with entities or terrain, not with each other
  - This means they also don't need a shape, so they take up less space
- Keep them in an old-fashioned C-style array
  - This limits the number of particles you can have (which is probably a good thing)
  - Keeps all the memory contiguous
  - Once you are trying to allocate more particles than you have room for, the oldest ones are kicked out first
- Tick them in your draw loop
  - Only having to iterate over them once
- Do them on the GPU (see CS 1230 lab)

Advanced Graphics

# DEFERRED LIGHTING

# Motivation



- Per-vertex lighting is ugly
  - Can increase number of vertices, but this requires lots of extra memory
- Per-pixel lighting looks a lot better
  - But is much slower than per-vertex, must calculate lighting equation for each **pixel** instead of each **vertex**.
  - Naïve implementation has many wasted calculations
    - Calculates lighting on pixels that are later overwritten by a closer triangle
- Deferred lighting removes most of the wasted calculations and provides further optimizations
- Deferred lighting is an optimization, not a lighting model. You still have to choose a lighting model (for example [Phong Lighting](#))

# Overview

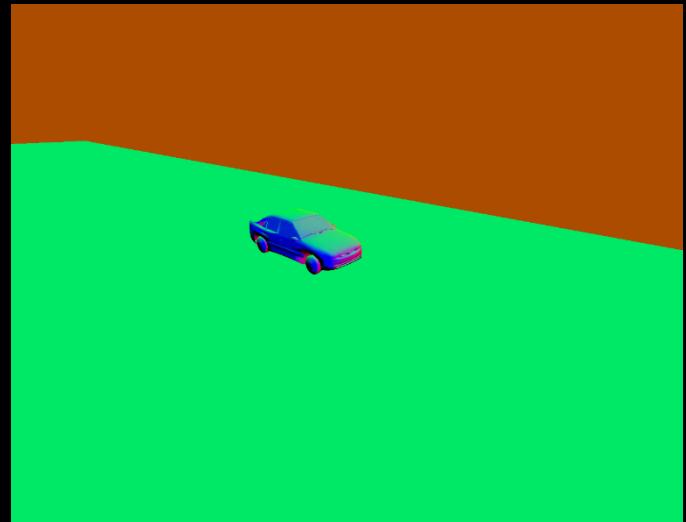
- How can we avoid wasted calculations?
  - Only calculate lighting once for each pixel
  - But the fragment shader has no way of knowing if the value it is calculating will be the final value for that pixel
  - Solution: Multiple passes
- First pass
  - Render geometry and keep track of data necessary to calculate lighting
- Second pass
  - Calculate diffuse and specular values that are independent of material
- Third Pass
  - Combine diffuse/specular from second pass with geometry and material (Object color for example) to complete lighting model

# Passes

- A “pass” just means generating a texture (or multiple textures).
- Use framebuffer objects (FBOs) to group textures
  - FBOs are basically a collection of textures
  - The FBO allows you to write to these textures (instead of writing to the screen)
  - Default framebuffer is the screen
    - `graphics->setDefaultFramebuffer();`
  - Look at FBO class methods
    - constructor
    - bind (sets framebuffer as active)
    - `getColorAttachment` (gets texture for one of the framebuffer’s color attachments)
  - And Graphics methods
    - `addFramebuffer`, `setFramebuffer`
- For example:
  - First pass writes to “Texture1” (using “FBO1”)
  - Second pass reads from “Texture1” and writes to “Texture2” (using “FBO2”)
  - Third pass reads from “Texture2” and writes to the screen

# First Pass

- Takes in all our geometry as input
  - Doesn't need material properties (ex. textures)
  - This just means you need to draw everything in the scene using your first pass shader
- It outputs exactly the information we need to calculate lighting
  - Normals
  - Positions
  - Shininess – store as alpha channel of normal



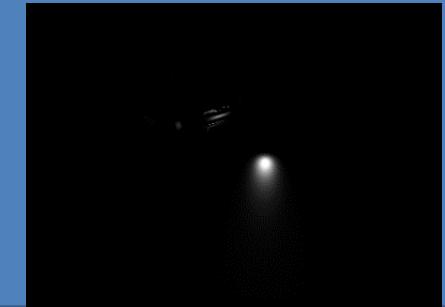
# Second Pass (1/2)

- Takes in normals, shininess and positions from first pass and light data
- Outputs diffuse and specular light contributions
  - Can save space by rendering to a single texture and storing specular contribution as alpha (but we only get monochromatic specular highlights)
  - Or render diffuse and specular to separate textures
- How do we send light data to GPU?
  - For each light:
    - Set necessary uniforms (position, direction, color, etc...)
    - Naïve: render full-screen quad to run the fragment shader on each pixel
      - Can do better, see slide 10
  - But each light would overwrite data from previous light
    - Solution: `graphics->enableBlendTest(Graphics::BLEND_FUNC::ADD)` for additive blending

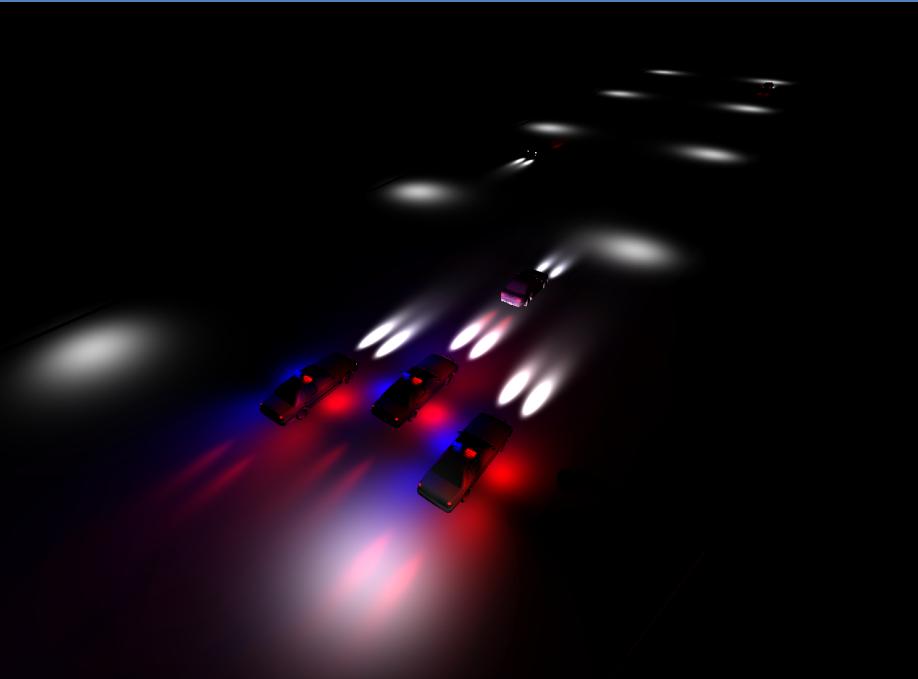
$$\text{Diffuse: } (\vec{n} \cdot \vec{l}) * \text{lightColor}$$



$$\text{Specular: } (\vec{r} \cdot \vec{v})^n * \text{lightColor}$$



# Second Pass (2/2)



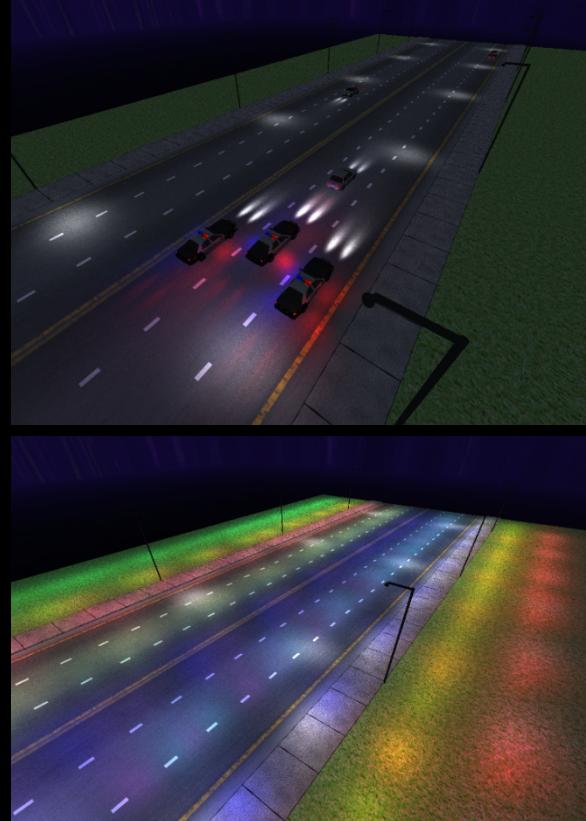
Diffuse



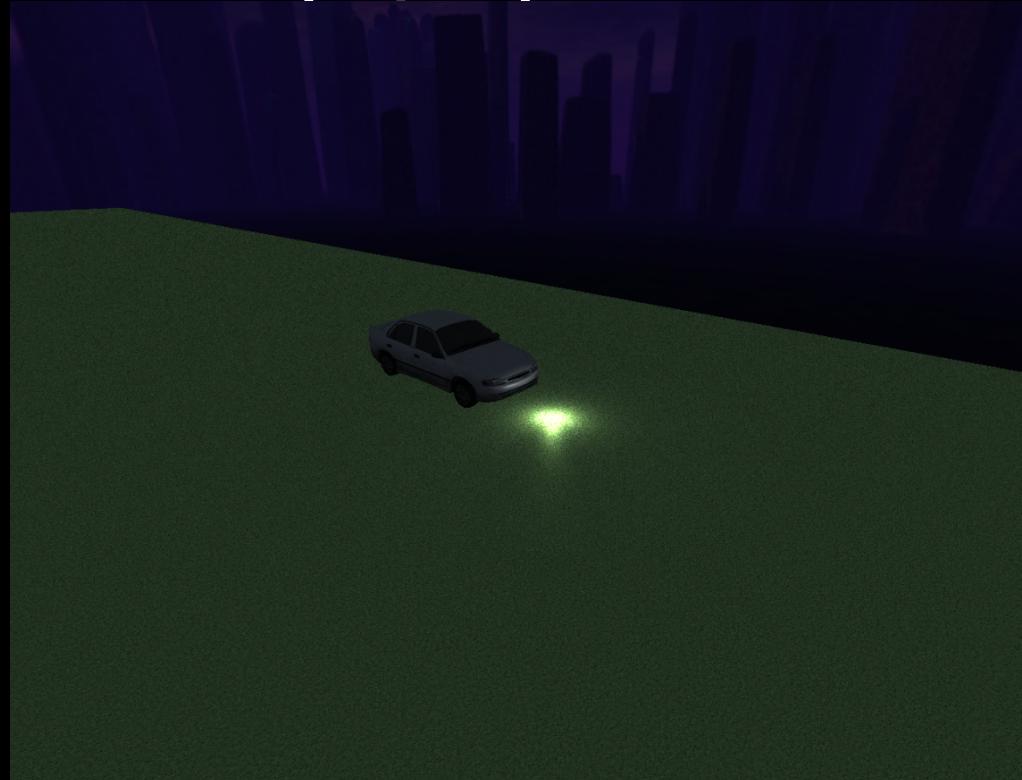
Specular

# Third Pass (1/2)

- Takes in diffuse and specular contribution from second pass and geometry, textures, etc... (whatever we need to calculate object's diffuse color)
- Render the scene again, this time applying any materials and finishing the lighting equation (i.e. finish calculating diffuse and specular term and add ambient + diffuse + specular)
- Output is our final lit scene which goes to the screen



# Third Pass (2/2)

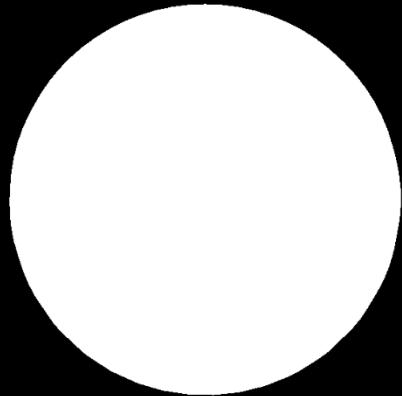


# Optimizations

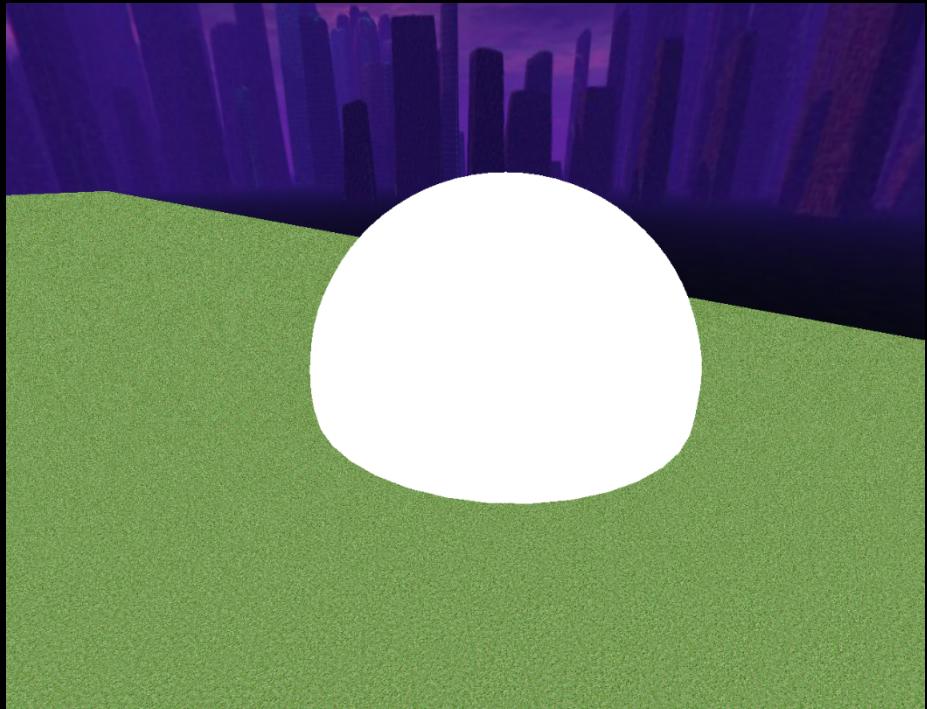
- Instead of calculating lighting on every pixel for every light, calculate lighting on only the subset of pixels that a light can possibly affect.
- Restrict the lighting calculations to the geometric shape that represents the volume of the light. In the second pass render this 3D shape instead of a full-screen quad.
  - Point light: sphere (radius usually based on attenuation)
  - Spot light: cone
  - Directional light: full-screen quad
- What if the camera is inside the shape of the light?
  - Represent light as a full-screen quad
- We will still have some wasted calculations, but this is much better (especially for small lights).

# Optimizations

How the light is represented



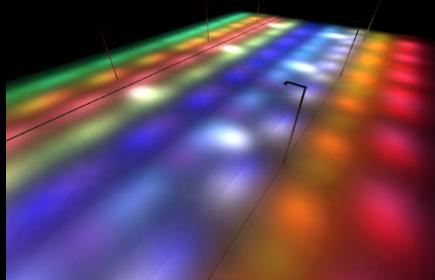
Light visualized in scene



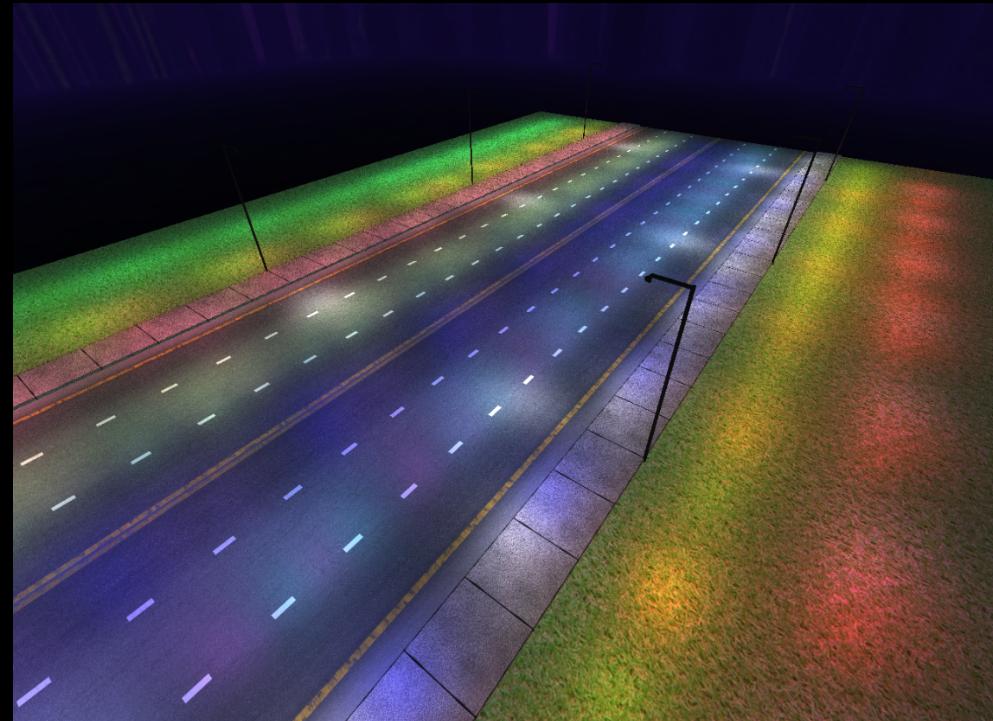
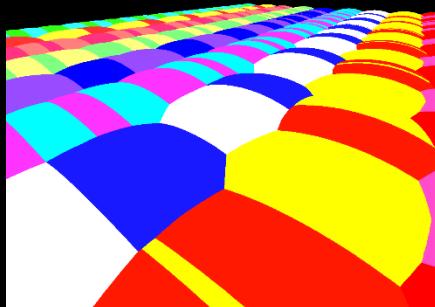
# Optimizations

Final scene

Diffuse contribution



Visualization of lights



# Optimizations

- The second pass needs to know the position of each pixel in world space
  - Our first pass shader can easily write this position to a texture
- Doing this uses an extra texture (i.e. twice as much memory)
- Instead, can use the depth buffer from the first pass.
  - First pass already uses a depth buffer, so we don't need any additional space.
- Depth buffer has z value from 0 to 1 for each pixel in screen space (convert x/y to screen space based on width/height).
- Use the (x,y,z) triplet in screen space and the inverse projection and view matrices to transform to world space.

# Deferred Shading

- Deferred shading is another method for speeding up per-pixel lighting, almost the same as deferred lighting.
  - Note that the word “shading” here doesn’t refer to interpolating lighting values, its just the name of this technique
- Uses only 2 passes
  - First pass renders geometry storing normals as in deferred lighting, but also stores all material properties (diffuse color, for example) in one or more additional textures
  - Second pass calculates lighting and uses material properties to calculate final pixel color
- Pros: Less computation (don’t need to render the scene twice)
- Cons: Uses more memory and bandwidth (passing extra textures around)

# Overview

- Deferred Lighting
  - 1. Render normals, positions, and shininess to textures
  - 2. Calculate diffuse and specular contributions for every light and output to textures
  - 3. Render scene again using diffuse/specular light data to calculate final pixel color (according to lighting model)
- Deferred Shading
  - 1. Render normals, positions, shininess, and material properties to textures
  - 2. Calculate lighting for every light and combine with material properties to output final pixel color

# Disadvantages of Deferred Rendering

- Can't easily handle transparency (this is a generic issue with z-buffer rendering techniques)
- Solutions:
  - Naïve: Sort transparent objects by distance
    - Slow
    - Can't handle intersecting transparent objects
  - Order-independent transparency: Depth peeling
  - Use forward-rendering for transparent objects
    - Forward-rendering is the standard rendering pipeline that you've been using
- Can't use traditional anti-aliasing techniques
  - MSAA (Multisample anti-aliasing), one of the most common AA techniques, doesn't work at all with deferred lighting
  - Usually use some sort of screen-space anti-aliasing instead (FXAA, TSAA, MLAA, SMAA)

Advanced Graphics

# VOLUMETRIC EFFECTS

# Volumetric Effects

- Volumetric glow (fake scattering)



# Volumetric Effects

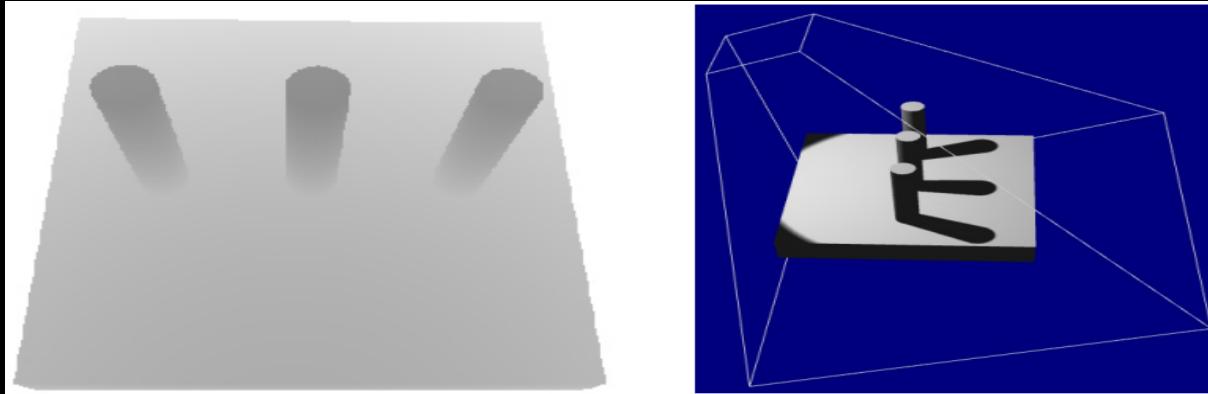
- Volumetric glow (fake scattering)
  - Blend glow color over every pixel
  - Fade off using closest distance from light source to line segment starting from eye and ending at object under pixel
  - Requires deferred shading for position of object
- Rendering to entire screen is expensive
  - Fade off to zero at some radius
  - Only need to draw pixels within that radius in world space, will be cheap for a far away effect
  - Render using inside-out sphere with that radius

Advanced Graphics

# SHADOW MAPPING

# Shadow Mapping

- Need to test whether a pixel is in shadow
  - Render scene from light's point of view
  - Depth map stores closest point to light
  - Render the scene from the camera, projecting each point back into the light's view frustum
  - Shadow if depth of projected point > depth map



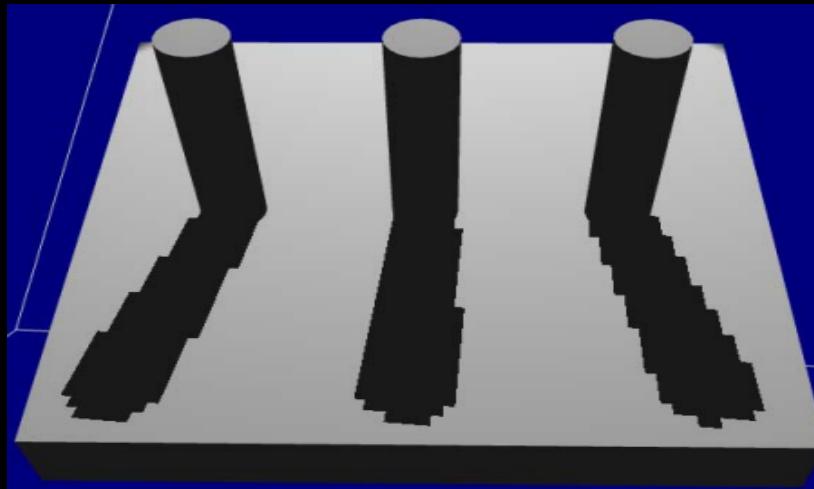
# Shadow Mapping

- Need to fit frustum to light
  - Directional light => parallel rays => orthographic
  - Spot light => frustum => perspective
  - Point light => rays in all directions => use cube map



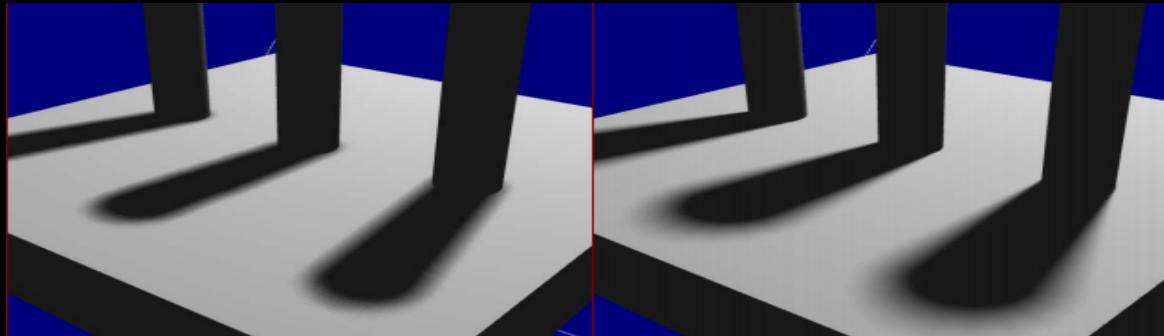
# Shadow Mapping

- Problem: Jagged edges
  - Shadow map resolution varies across scene
  - Increasing resolution helps, but uses more memory



# Shadow Mapping

- Fix: blur or fuzz out boundaries
  - Multiple nearby shadow tests are made per pixel and are averaged together
  - Called PCF: Percentage Closer Filtering
    - May use randomized sample patterns
    - May use variable blur size since shadows get more blurry away from caster and area lights



# Shadow Mapping

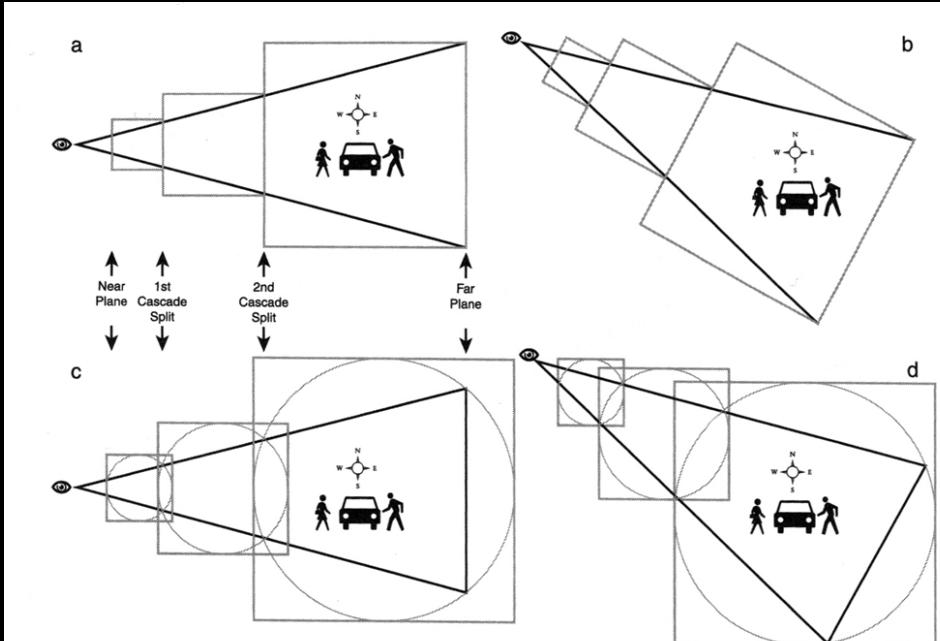
- Fix: Average out over multiple frames
  - Reproject previous frame (if moving camera)
  - Jitter shadow map per frame for more samples
  - Weight by confidence (distance to texel center)
  - <http://levelofdetail.wordpress.com/2008/11/06/pixel-correct-shadow-maps-with-temporal-reprojection/>



# Cascaded Shadow Maps

- Shadow mapping has problems
  - Resolution varies across scene
  - One shadow map per object doesn't scale
- Idea: fit several shadow maps to camera
  - Want uniform shadow map density in screen-space
  - Objects near eye require higher world-space density than objects far away
  - Use a cascade of 4 or 5 shadow maps
  - Each one is for a certain depth range of the scene
- Used in almost all modern games

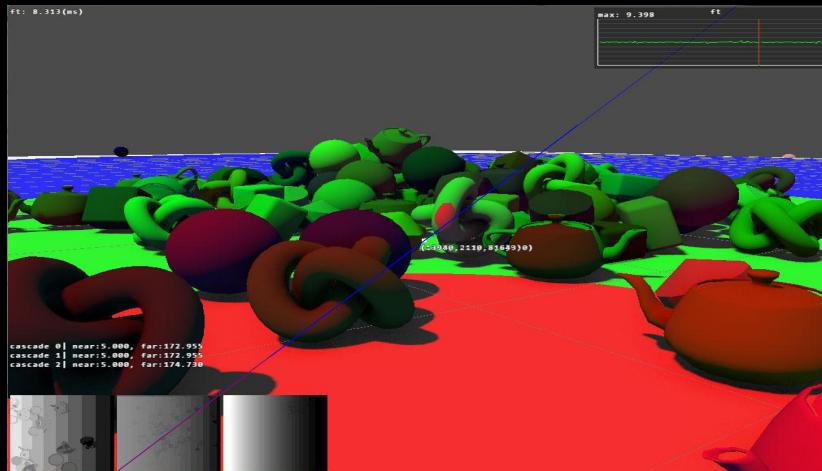
# Cascaded Shadow Maps



**FIGURE 4.1.2** The view frustum in world space split into three cascade frustums and their corresponding shadow map coverage. We use a top view with the light direction pointing straight down the horizontal world plane.

# Cascaded Shadow Maps

- Use depth in shader to choose cascade
  - Can blend between two closest cascades to smooth out discontinuities



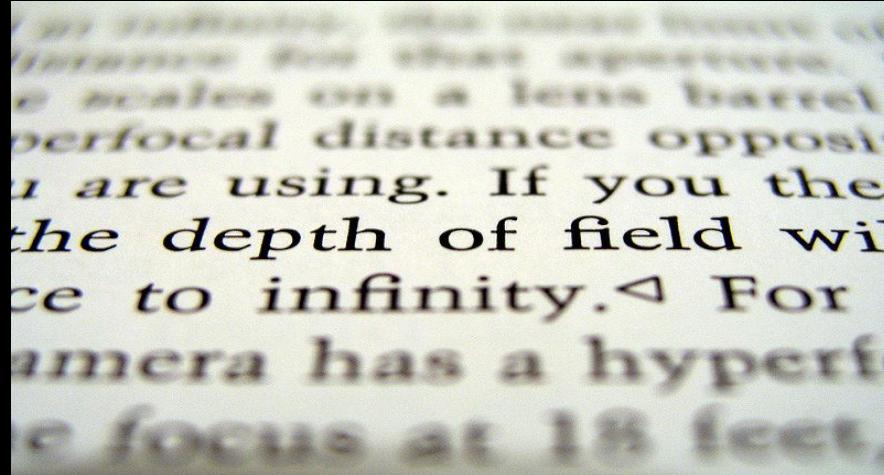
Scene with a cascade of 3

Advanced Graphics

# DEPTH OF FIELD

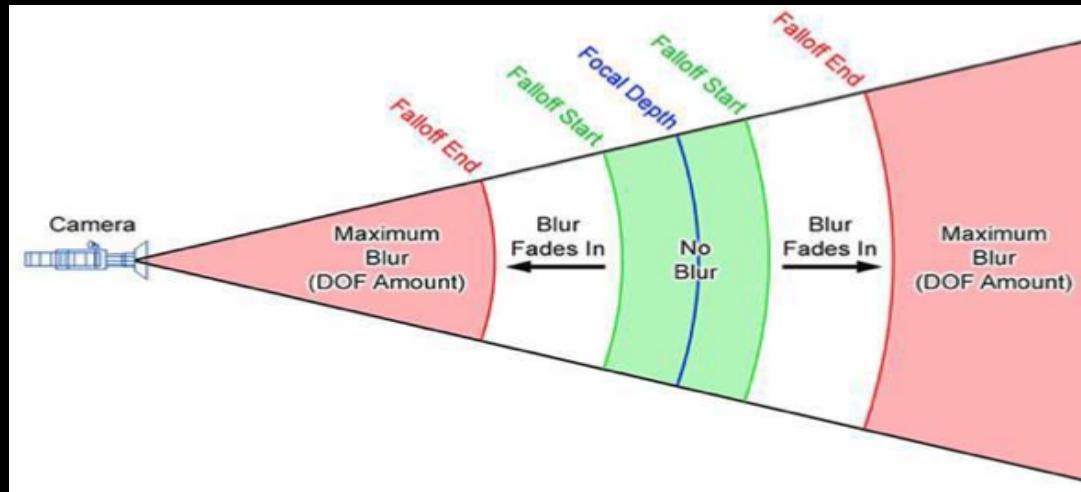
# Depth of Field

- Out of focus blur in real cameras
  - Only one depth where objects are in focus
  - Focal blur increases in both directions away from that depth



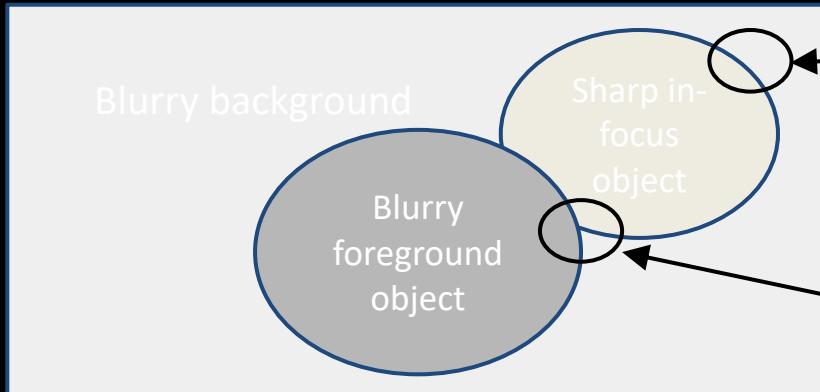
# Depth of Field

- Model with post-process blur
  - Vary blur radius based on scene depth: slow
  - Interpolate between 3 images blurred with different radii: fast



# Depth of Field

- Discontinuities are problematic (halos)
  - Don't use sharp objects in background blur
  - Blur over sharp objects for foreground objects



Background blur  
should not use in-  
focus pixels

Blurry foreground objects  
should contribute to sharp  
objects

# Depth of Field in Starcraft II

- Avoid sharp halos
  - Buffer of per-pixel blur radius
  - Weigh blur samples by radius buffer at sample point
  - Renormalize to sum to 1 again
- Halos around blurry objects
  - Compute blurred radius buffer
  - Compute blurred depth buffer (local average depth)
  - If average depth < current depth, use radius from blurred buffer, otherwise use radius from sharp buffer



DOF in Starcraft II  
cutscene

Advanced Graphics

# LIGHTING EFFECTS

# Screen-Space Ambient Occlusion

- Darken ambient term in occluded areas
  - Approximates indirect lighting

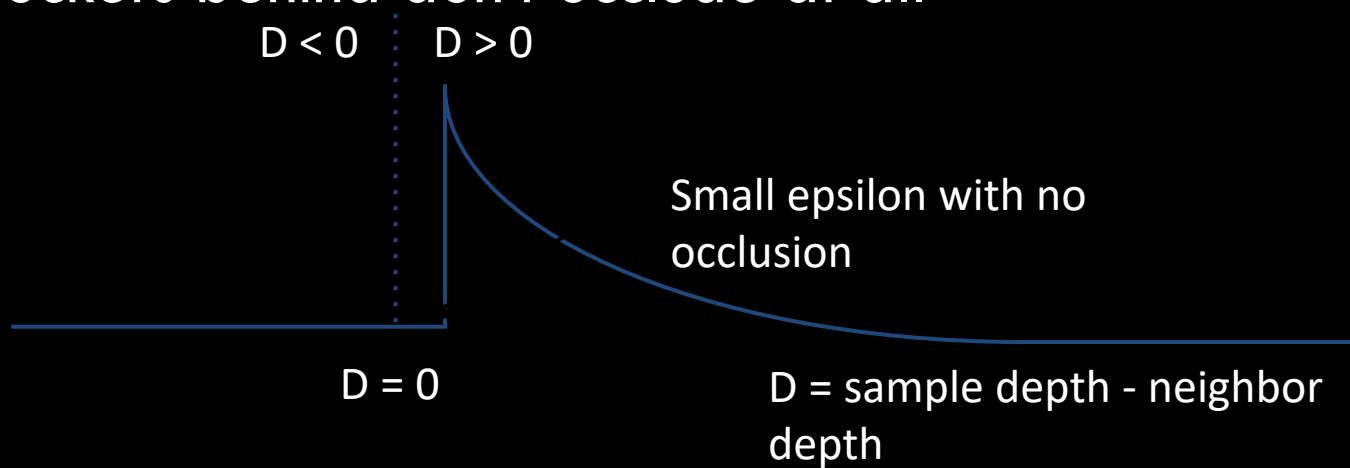


# Screen-Space Ambient Occlusion

- Calculating occlusion
  - Probe nearby geometry using raycasting
  - Shoot rays in a hemisphere out of surface
  - Objects in close proximity cause darkening
- Idea: per-pixel occlusion approximation
  - Flatten raycasting to 2D in the image plane
  - Sample the depth of 8 to 32 neighboring pixels (requires deferred shading)
  - Don't count off-image samples as occlusions
  - Compare neighbor depth to 3D sample depth
  - If neighbor is in front, occlusion may be occurring

# Screen-Space Ambient Occlusion

- Occlusion falloff function
  - Blockers closer to the sample should occlude more
  - Blockers far from the sample don't occlude at all
  - Blockers behind don't occlude at all



# Real Time Local Reflections

- Used in Crysis 2



# Real Time Local Reflections

- Used in Crysis 2



# Real Time Local Reflections

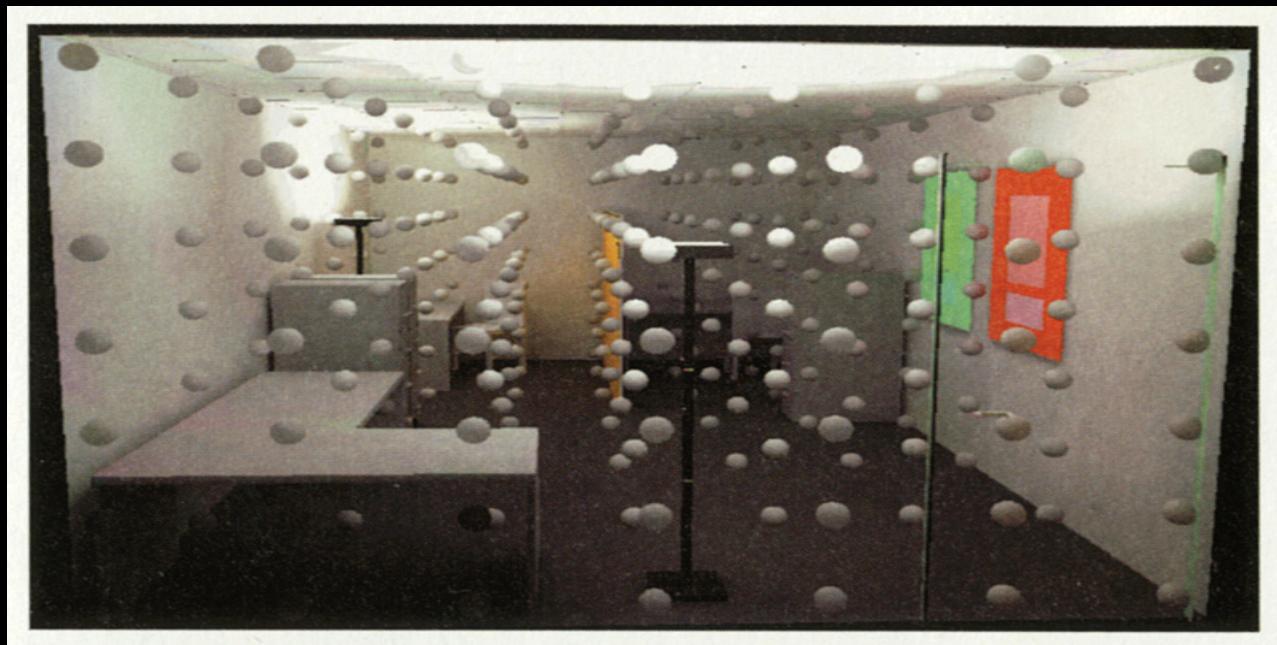
- Existing reflection methods (rasterization)
  - Render scene flipped about a plane
    - Use rendered scene as reflection
    - Only works for planar surfaces
  - Render scene from a point of view into cube map
    - Look up into cube map using reflection vector
    - Only works for small objects
- Reflections are expensive with rasterization
  - Need to render scene once per planar surface or per reflective object
  - Raytracing is much more straightforward

# Real Time Local Reflections

- Raytrace reflections in screen space
  - Compute reflection vector per pixel using depth and normal from G-buffer
  - Raymarch along reflection vector
  - Project ray into 2D and check if scene depth is within threshold of ray depth
  - If so, use color from previous frame as reflection
- Edge cases (no data)
  - Fade out as reflection faces the camera
  - Fade out reflections off screen edge

# Real Time Global Illumination

- Precomputed with irradiance cache

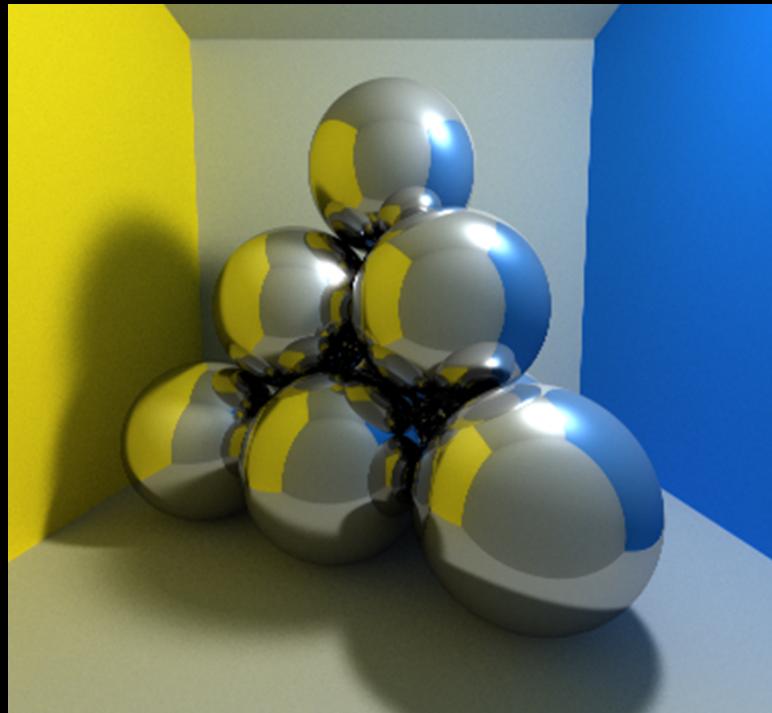


# Real Time Global Illumination

- Precomputed with irradiance cache
  - Lightmaps for static objects
  - 3D grid of irradiance samples for dynamic objects
    - Each sample is snapshot of all light coming into a point
    - Think cube map, usually compressed using spherical harmonics
- Animated lightmaps
  - Static scene with moving light restricted to a path
- Precomputed Radiance Transfer (PRT)
  - Lightmap that can be queried by incident light angle
  - Lighting solution stored compressed using SH

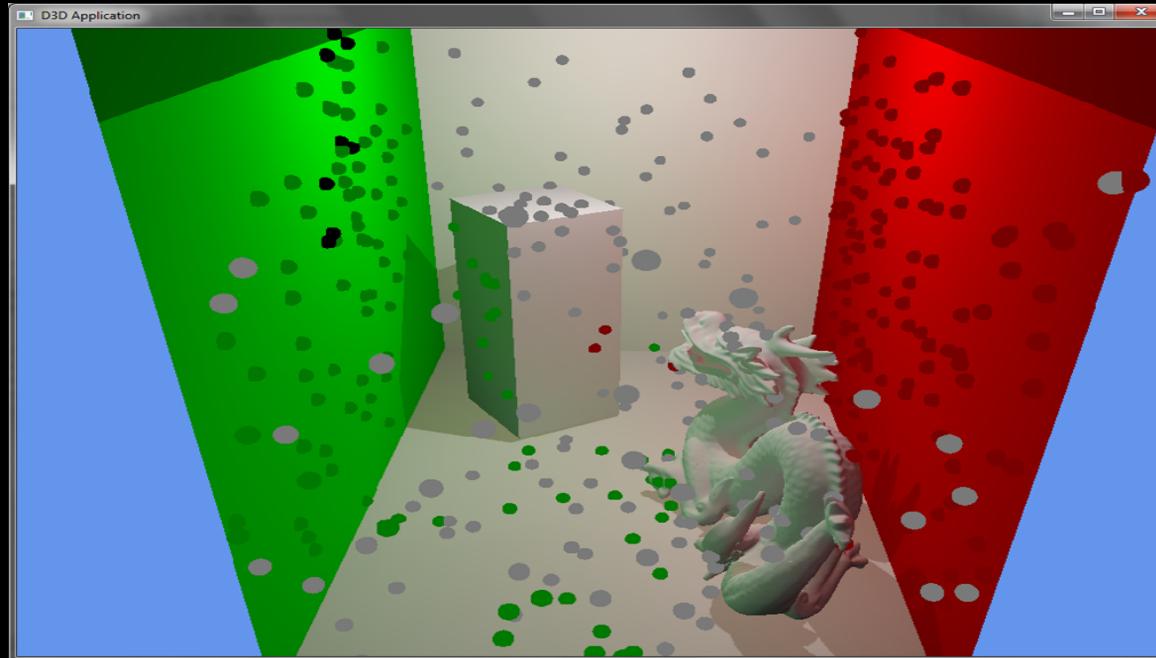
# Real Time Global Illumination

- Path tracing directly
  - Not used much in games,  
technology still  
advancing



# Real Time Global Illumination

- Instant radiosity



# Real Time Global Illumination

- Instant radiosity
  - Shoot some photons into the scene ( $\sim 200$ )
  - Only do one bounce
  - Virtual Point Light (VPL) where they land
- Rendering
  - Direct lighting: as usual (shadow maps)
  - Indirect lighting: each VPL becomes point light
- Updating
  - Randomized nature means lots of noise
  - Cache valid VPLs between frames

# CLASS 8

C++ Tip of the Week

C++ Tip of the Week

# LAMBDAS

# Lambdas

- **Inline functors:**

```
int i = ...;
double terrainHeight = ...;
auto adjustEntityHeight = [terrainHeight, i](Vector3 &pos) {
    pos.y = terrainHeight*complicatedMath();
};

... // code that would use this math multiple times

// template
auto functionName = [capturedVariables](input arguments) { ... };
```

# Lambda syntax

```
auto functionName = [capturedVariables](input arguments) { ... };
```

- **auto** tells the compiler to determine the type
  - in this case, it's a generated lambda definition
  - **auto** only works like this starting in C++11
    - Before that it declared a variable as ‘local’ storage, and now is almost never used
- **Captured variables give the function access to variables in the surrounding scope**

# More lambda syntax

- To capture member variables, pass `this` as a captured variable

```
auto func = [this](int intput) { ... };
```

- You can optionally specify the return type for clarity

```
auto functionName = [capturedVariables](input arguments) ->  
returnType { ... };
```

# Even more lambda syntax

```
auto func = [&localOne, localTwo](int intput) { ... };
```

- localOne is passed by reference, whereas localTwo is passed as a copy (by default)
- Use '=' or '&' to capture all local variables in scope

```
auto func = [&](int intput) { ... }; // captures all variables by reference
```

```
auto func = [=, &localTwo](int intput) { ... }; // captures all variables // by copy, except localTwo which is captured by reference
```

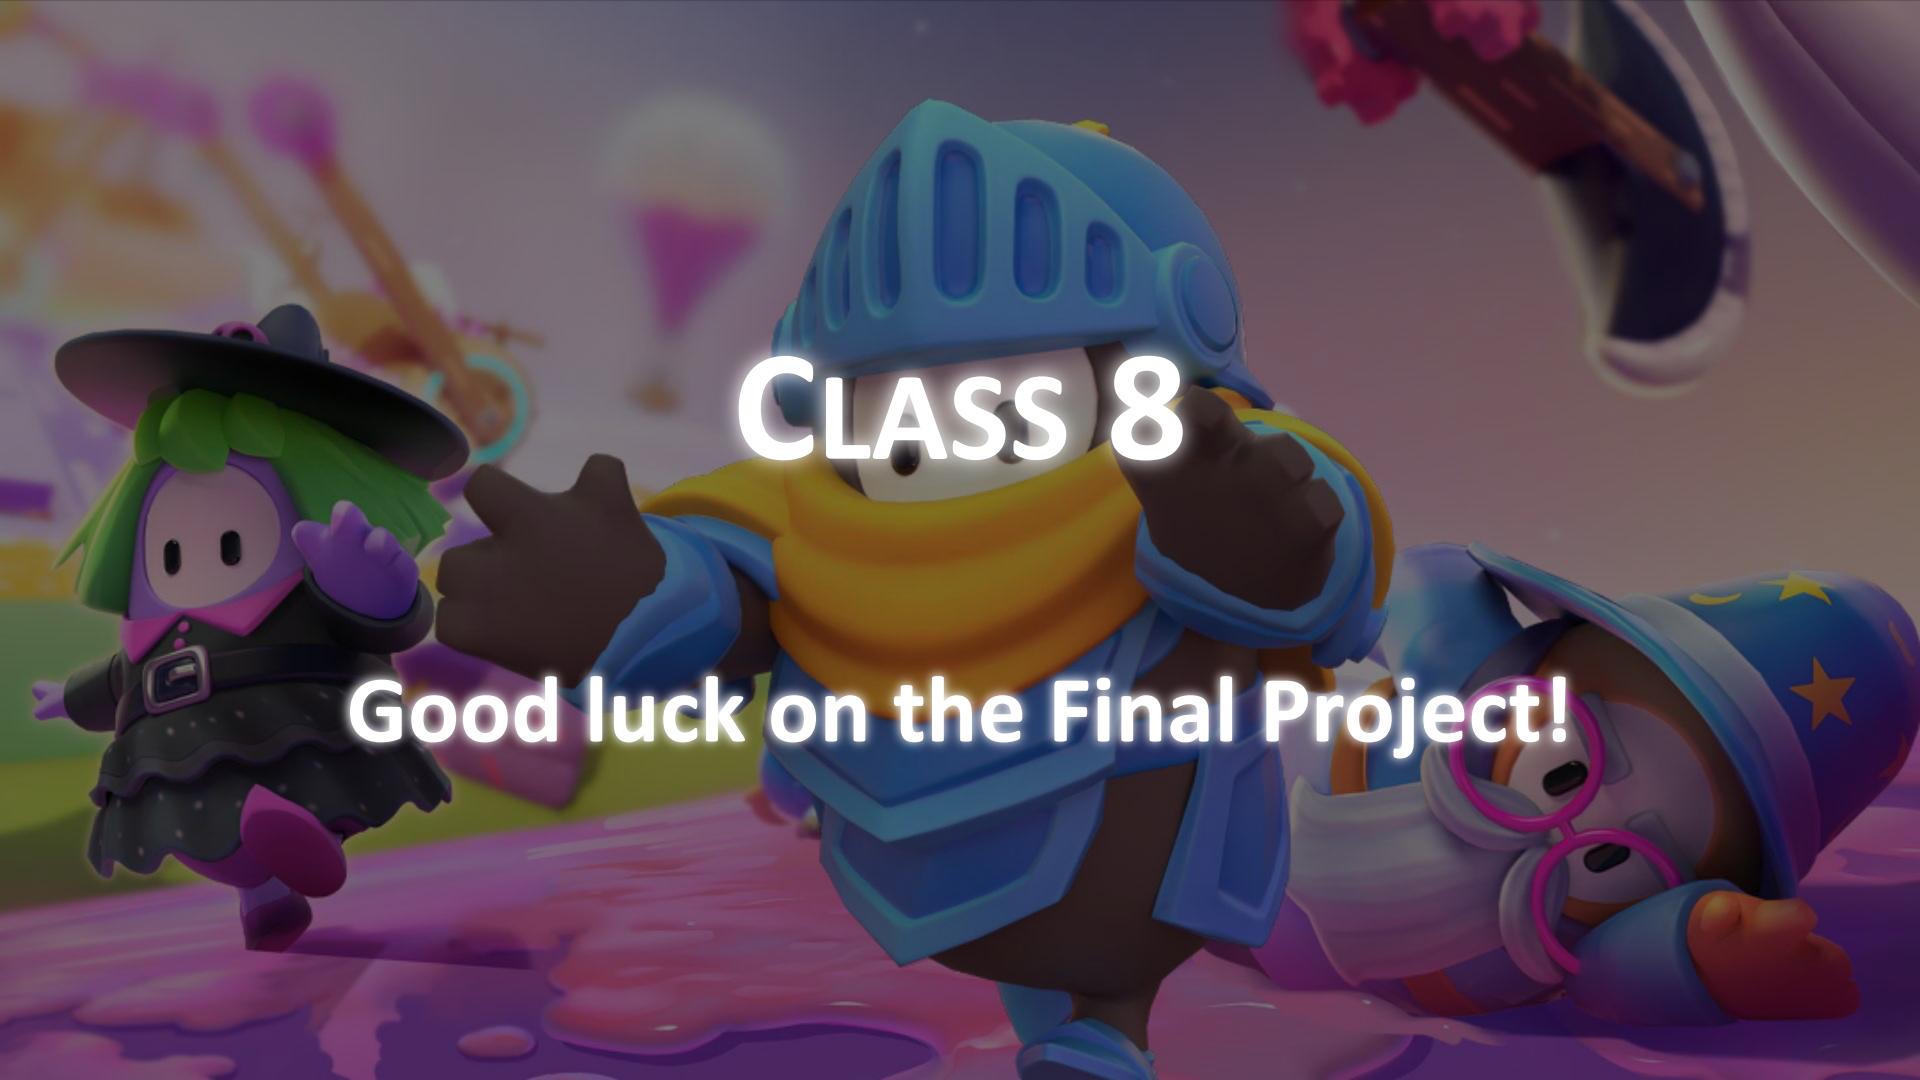
# Using Lambda anonymously

- As you probably guessed from the name `lambda`, you can also use them completely anonymously:

```
std::vector<int> some_list{ 1, 2, 3, 4, 5 };
int total = 0;
std::for_each(begin(some_list), end(some_list), [&total](int x)
{
    total += x;
});
```

# Lambda based algorithms

- The standard library contains a bunch of helpful functions that make use of lamdas and lists
- Besides `for_each`, some examples are `transform`, `count_if`, `remove_if`, and `binary_search`
- Read more about lambdas:  
<https://stackoverflow.com/questions/7627098/what-is-a-lambda-expression-in-c11>



# CLASS 8

Good luck on the Final Project!